

# GITting started with reproducibility: An introduction to `git` and `knitr`

## Biostatistics Student Association Computing Workshop

Nick Seewald

Department of Statistics  
University of Michigan

January 29, 2016

# Why do I care about reproducibility?

Reproducible research is a hallmark of the scientific method, but we're pretty bad at it.

*In 2012, a researcher then at the biotechnology company Amgen wrote in Nature that when his team tried to reproduce 53 landmark cancer studies, they could replicate just six. And according to a news report in Nature, a project aiming to reproduce the findings of 100 psychology papers has managed to replicate results for only 39 of them (the project's findings are still under peer review).*

"What Science Can Tell Us About Bad Science", *The Atlantic*, September 2015. <http://www.theatlantic.com/magazine/archive/2015/09/a-scientific-look-at-bad-science/399371/>

# But I'm a Biostatistician!

- ▶ Reproducibility is important in both science AND statistics!
- ▶ As statisticians, we need to be able to reproduce our results on the same data set
  - ▶ This means we have to write reports in a way that minimizes error and write code so that we can get the same results years later.

# Agenda

1. Git: A “version control” tool used for collaborating and maintaining different versions of a file, typically for code.
  - ▶ Great for collaborating, or just saving your own ass.
  - ▶ Often used in conjunction with *GitHub*, an online repository storage service.
2. knitr: An R package that lets you create documents containing R code and output.
  - ▶ Keep everything you need to generate a report (e.g., for research, homework, or 699) in one place!
  - ▶ My favorite part: Update code without having to re-create tables! (This is where errors creep in!)

# A brief warning



Source: <https://xkcd.com/1597/>

# Setup

Create a GitHub account, then download either Git or GitHub Desktop.

- ▶ Pure Git (i.e., just command-line tools):  
<https://git-scm.com/downloads>
- ▶ GitHub Desktop (GUI & command-line tools):  
<https://desktop.github.com/>

# Forking a Repository

A *fork* is a copy of a repository. Forks let you

- ▶ make changes to an existing repo without affecting the original project
- ▶ use the existing project as a starting point for your own

One of the benefits of open source!

**Exercise:** Fork my bsa-computing repository.

# Cloning a Repository

- ▶ To create a local copy of an existing git repository, use `git clone [url] [directory-name]`.
- ▶ In a terminal (Mac, Linux) or the Git Shell (Windows), navigate to the folder you want to clone the repository into.
- ▶ **Exercise:** Clone your forked bsa-computing repository onto your computer. The URL will be of the form `https://github.com/[username]/bsa-computing.git`



# Make Changes!



- ▶ You now have a copy of both the current version of `bsa-computing` *and* access to every previous version.
- ▶ This clone is NOT automatically synced, à la Dropbox
  - ▶ Anything you break is completely isolated from the pristine copy on GitHub and your previous “commits”.
- ▶ **Exercise:** Add to `food-exercise.md`, and save your changes.

# Making Commits and Pushing

- ▶ Once you've accomplished a relatively small, but still significant task, you'll want to “commit” your code to the repository.
- ▶ This creates a labeled snapshot of the directory at the time of the commit.
- ▶ Syncing your commits with the existing remote repository is called “pushing”.

# Shell Commands for Commits and Push/Pull

- ▶ To *pull* down the most recent version of the remote repository, use `git pull`
- ▶ To make and push commits
  - ▶ Check what you've changed with `git status` and `git diff`.
  - ▶ Add files to the commit using `git add [files]`.
  - ▶ Create a commit using `git commit -m [message]`
  - ▶ Push your commit to the remote repository using `git push`

# Merging Your Fork with the Original

On GitHub, create a *pull request* to merge your changes into my repository. The owner (me) will be able to look at and accept/decline them.

# What is knitr?

- ▶ knitr lets you embed code and output from R into  $\text{\LaTeX}$ , HTML, RMarkdown, etc. I'll be using  $\text{\LaTeX}$ .
- ▶ Might require a bit more thought with R code to get the correct output, but it's worth it.

The knitr Bible: <http://yihui.name/knitr/>

# Chunks

knitr separates R code from  $\text{\LaTeX}$  using structures called “chunks”.

## Chunk Syntax

```
<<list(chunkname, eval=FALSE)>>= @
```

# Chunk Options


When you define a chunk, you also can include a number of options that tell `knitr` how to interpret the R code inside. Some basic important options are:

- ▶ `echo` (`TRUE`): Whether or not to include the R code itself in the document.
- ▶ `results`: How to display results. Default is in a `LaTeX verbatim` environment.
  - ▶ `asis`: Write raw results into the document (e.g., with `xtable`)
- ▶ `include` (`TRUE`): Whether to include the chunk output in the document. The code is still run. Useful for data steps.
- ▶ `cache` (`FALSE`): Stashes the output of a chunk to avoid running it every time you compile.

More options and more details at <http://yihui.name/knitr/options/>.

# Example: Figure

**Fitted Gaussian Model for Height and Weight in Males**





## Example: Tables

There are 137 women in this dataset, out of 210 total individuals.

	Mean	Std. Dev.	Median
Height	67.36	4.46	67.00
Weight	145.66	32.69	140.00

Table: Summary Statistics for All Individuals

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1	65.00	70.00	72.00	71.66	73.00	79.00
2	125.00	155.00	170.00	175.60	185.00	280.00

Table: Summary Statistics for Men